

Lifelong Learning Techniques

for an Origami Robot

by

E. De Vroey

Student number: 4685156

Contents

1	Introduction	1
1.1	Application to soft (origami) robots	1
1.2	Challenges	1
2	Background and definition	3
3	Blackbox methods	5
3.1	Learning control policies	5
3.2	Learning the dynamic model	5
4	Learning architectures	7
4.1	Dual architectures	7
4.2	Architectural expansions	7
5	Storage and retrieval	9
6	Discussion and conclusion	11

1

Introduction

1.1. Application to soft (origami) robots

In the field of soft robotics, lifelong learning are particularly interesting.

1.2. Challenges

(Kim et al., 2021; Lesort et al., 2020).

How can lifelong learning methods be employed to manage the evolving dynamics due to wear and tear in the control of a soft origami robot over its operational lifespan?

2

Background and definition

(Annaswamy & Fradkov, 2021; Landau et al., 2011; Lesort et al., 2020).

Blackbox methods

(Alessi et al., 2023; Gijssberts & Metta, 2011; Nguyen-Tuong et al., 2009; Romeres et al., 2016; Thuruthel et al., 2019; Zhou et al., 2022).

3.1. Learning control policies

3.2. Learning the dynamic model

One way of having continually adapting control for soft robots, is by implementing a (traditional) control scheme that uses a dynamic model (either forward or inverse) of the robot and then to incrementally update that model. The idea of learning a dynamic model for soft robots has advantages even outside the context of lifelong adaptation: many soft robots have such a large amount of degrees of freedom and such complex and nonlinear dynamics, that analytically deriving a dynamic model is both cumbersome and often insufficiently accurate (Gillespie et al., 2018; Giorrelli et al., 2015; Kim et al., 2021; Thuruthel et al., 2019).


Blackbox methods also allow for learning both forward and inverse dynamic models and can thus be used in several control schemes. When learning a forward model, the robot's actions and resulting states can be used as the inputs and outputs, respectively, this is so-called *direct modeling* (Nguyen-Tuong & Peters, 2011b). In the case of learning inverse dynamic models, such direct approaches might not always be applicable, due to the multiple solutions that exist for the inverse dynamics of robots with redundant degrees of freedom. Instead, an *indirect-modeling* approach can be taken, where the model is trained to minimize the error of the feedback controller (Nguyen-Tuong & Peters, 2011b).

The form that the dynamic model takes can vary. In (Gillespie et al., 2018), a neural network is trained to predict the next state of the robot, given the current state and a control action, the analytical gradients of these outputs w.r.t. to the inputs are made available during the process of backpropagation and can be used in matrices A and B to form a more interpretable state-space system. In (Thuruthel et al., 2019) on the other hand, a forward model is learned for open-loop control, using the control action, the current state, and the previous state as input, and the next state as output to train a recurrent neural network. This forward model is kept in its more abstract blackbox form, and is directly used to generate trajectories in open-loop control.

refine structure

Likewise, since the lifelong adapting nature of the techniques is embedded in the dynamic model, the choice of control scheme in which the learned dynamic model is used can vary from implementation to implementation. The aforementioned matrices A and B from (Gillespie et al., 2018) are used in a model predictive control setting, while the generated trajectories from (Thuruthel et al., 2019) are used as samples to train a blackbox control policy (the dy-

dynamic model is in essence an intermediary step in order to create a blackbox policy as in section 3.1).



Learning architectures

4.1. Dual architectures

(Lu et al., 2019, 2020; Mallya et al., 2018; Schwarz et al., 2018; Shin et al., 2017; Sprechmann et al., 2018).

4.2. Architectural expansions

(Chatzilygeroudis et al., 2016; Gillespie et al., 2018; Nazeer et al., 2023).

5

Storage and retrieval

(Aljundi et al., 2016; Mouret & Clune, 2015; Nguyen-Tuong & Peters, 2011a; Schulman et al., 2017; Shin et al., 2017; Sprechmann et al., 2018; Wang et al., 2022; Xie & Finn, 2021).

6

Discussion and conclusion

Bibliography

- Alessi, C., Hauser, H., Lucantonio, A., & Falotico, E. (2023). Learning a controller for soft robotic arms and testing its generalization to new observations, dynamics, and tasks. *2023 IEEE International Conference on Soft Robotics (RoboSoft)*, 1–7. <https://doi.org/10.1109/RoboSoft55895.2023.10121988>
- Aljundi, R., Chakravarty, P., & Tuytelaars, T. (2016). Expert gate: Lifelong learning with a network of experts. <http://arxiv.org/abs/1611.06194>
- Annaswamy, A. M., & Fradkov, A. L. (2021, January). A historical perspective of adaptive control and learning. <https://doi.org/10.1016/j.arcontrol.2021.10.014>
- Chatzilygeroudis, K., Cully, A., & Mouret, J.-B. (2016). Towards semi-episodic learning for robot damage recovery. <http://arxiv.org/abs/1610.01407>
- Gijsberts, A., & Metta, G. (2011). Incremental learning of robot dynamics using random features. *2011 IEEE International Conference on Robotics and Automation*, 951–956. <https://doi.org/10.1109/ICRA.2011.5980191>
- Gillespie, M. T., Best, C. M., Townsend, E. C., Wingate, D., & Killpack, M. D. (2018). Learning nonlinear dynamic models of soft robots for model predictive control with neural networks. *2018 IEEE International Conference on Soft Robotics (RoboSoft)*, 39–45. <https://doi.org/10.1109/ROBOSOFT.2018.8404894>
- Giorelli, M., Renda, F., Calisti, M., Arienti, A., Ferri, G., & Laschi, C. (2015). Learning the inverse kinetics of an octopus-like manipulator in three-dimensional space. *Bioinspiration & Biomimetics*, 10, 035006. <https://doi.org/10.1088/1748-3190/10/3/035006>
- Kim, D., Kim, S.-H., Kim, T., Kang, B. B., Lee, M., Park, W., Ku, S., Kim, D., Kwon, J., Lee, H., Bae, J., Park, Y.-L., Cho, K.-J., & Jo, S. (2021). Review of machine learning methods in soft robotics (G. Gu, Ed.). *PLOS ONE*, 16, e0246102. <https://doi.org/10.1371/journal.pone.0246102>
- Landau, I. D., Lozano, R., M'Saad, M., & Karimi, A. (2011). *Adaptive control*. Springer London. <https://doi.org/10.1007/978-0-85729-664-1>
- Lesort, T., Lomonaco, V., Stoian, A., Maltoni, D., Filliat, D., & Díaz-Rodríguez, N. (2020). Continual learning for robotics: Definition, framework, learning strategies, opportunities and challenges. *Information Fusion*, 58, 52–68. <https://doi.org/10.1016/j.inffus.2019.12.004>
- Lu, K., Grover, A., Abbeel, P., & Mordatch, I. (2020). Reset-free lifelong learning with skill-space planning. <http://arxiv.org/abs/2012.03548>
- Lu, K., Mordatch, I., & Abbeel, P. (2019). Adaptive online planning for continual lifelong learning. <http://arxiv.org/abs/1912.01188>
- Mallya, A., Davis, D., & Lazebnik, S. (2018). Piggyback: Adapting a single network to multiple tasks by learning to mask weights. <http://arxiv.org/abs/1801.06519>
- Mouret, J.-B., & Clune, J. (2015). Illuminating search spaces by mapping elites. <http://arxiv.org/abs/1504.04909>
- Nazeer, M. S., Bianchi, D., Campinoti, G., Laschi, C., & Falotico, E. (2023). Policy adaptation using an online regressing network in a soft robotic arm. *2023 IEEE International Conference on Soft Robotics (RoboSoft)*, 1–7. <https://doi.org/10.1109/RoboSoft55895.2023.10121927>

- Nguyen-Tuong, D., & Peters, J. (2011a). Incremental online sparsification for model learning in real-time robot control. *Neurocomputing*, 74, 1859–1867. <https://doi.org/10.1016/j.neucom.2010.06.033>
- Nguyen-Tuong, D., & Peters, J. (2011b). Model learning for robot control: A survey. *Cognitive Processing*, 12, 319–340. <https://doi.org/10.1007/s10339-011-0404-1>
- Nguyen-Tuong, D., Seeger, M., & Peters, J. (2009). Model learning with local gaussian process regression. *Advanced Robotics*, 23, 2015–2034. <https://doi.org/10.1163/016918609X12529286896>
- Romeres, D., Zorzi, M., Camoriano, R., & Chiuso, A. (2016). Online semi-parametric learning for inverse dynamics modeling. <http://arxiv.org/abs/1603.05412>
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. <http://arxiv.org/abs/1707.06347>
- Schwarz, J., Luketina, J., Czarnecki, W. M., Grabska-Barwinska, A., Teh, Y. W., Pascanu, R., & Hadsell, R. (2018). Progress & compress: A scalable framework for continual learning. <http://arxiv.org/abs/1805.06370>
- Shin, H., Lee, J. K., Kim, J., & Kim, J. (2017). Continual learning with deep generative replay. <http://arxiv.org/abs/1705.08690>
- Sprechmann, P., Jayakumar, S. M., Rae, J. W., Pritzel, A., Badia, A. P., Uria, B., Vinyals, O., Hassabis, D., Pascanu, R., & Blundell, C. (2018). Memory-based parameter adaptation. <http://arxiv.org/abs/1802.10542>
- Thuruthel, T. G., Falotico, E., Renda, F., & Laschi, C. (2019). Model-based reinforcement learning for closed-loop dynamic control of soft robotic manipulators. *IEEE Transactions on Robotics*, 35, 124–134. <https://doi.org/10.1109/TRO.2018.2878318>
- Wang, Z., Chen, C., & Dong, D. (2022). Lifelong incremental reinforcement learning with online bayesian inference. *IEEE Transactions on Neural Networks and Learning Systems*, 33, 4003–4016. <https://doi.org/10.1109/TNNLS.2021.3055499>
- Xie, A., & Finn, C. (2021). Lifelong robotic reinforcement learning by retaining experiences. <http://arxiv.org/abs/2109.09180>
- Zhou, W., Bohez, S., Humplik, J., Abdolmaleki, A., Rao, D., Wulfmeier, M., Haarnoja, T., & Heess, N. (2022). Forgetting and imbalance in robot lifelong learning with off-policy data. <http://arxiv.org/abs/2204.05893>