

Lifelong Learning Techniques

for an Origami Robot

by

E. De Vroey

Student number: 4685156

Contents

1	Introduction	1
1.1	Application to soft (origami) robots	1
1.2	Challenges	1
2	Background and definition	3
3	Blackbox methods	5
3.1	Learning control policies	5
3.2	Learning the dynamic model	6
3.3	Incremental learning.	6
3.4	Summary and conclusions	7
4	Learning architectures	9
4.1	Dual architectures	9
4.2	Architectural expansions	9
5	Storage and retrieval	11
6	Discussion and conclusion	13

1

Introduction

1.1. Application to soft (origami) robots

In the field of soft robotics, lifelong learning are particularly interesting.

1.2. Challenges

(Kim et al., 2021; Lesort et al., 2020).

How can lifelong learning methods be employed to manage the evolving dynamics due to wear and tear in the control of a soft origami robot over its operational lifespan?

2

Background and definition

(Annaswamy & Fradkov, 2021; Landau et al., 2011; Lesort et al., 2020).

3

Blackbox methods

general intro, define "blackbox"

3.1. Learning control policies

A perfect example of a blackbox method is given by Alessi et al. (2023). In their study they used Proximal Policy Optimization (PPO) (Schulman et al., 2017) to obtain a feedforward neural network (FNN) that can generalize over new dynamics and tasks. The controller is a blackbox: a desired trajectory, several measurements of the robots state, and other values of interest, such as the error of robot's end-effector tip, are used as inputs, resulting in a control action. It essentially acts in a manner akin to a feedback controller using an inverse dynamics model, only in this case it is unknown what is happening *under the hood*.

The previous example is "*perfect*" in the sense that it perfectly embodies the simplistic nature of blackbox methods, however it is worth noting that other techniques, such as more complicated architectures or combinations of multiple models, can be used in combination with these blackbox methods. Such architectural modifications will be discussed in chapter 4.

Apart from Alessi et al. (2023), several other examples of blackbox policies can be found in the literature.

In Zhou et al. (2022), two policies are learned for the control of a small humanoid robot. Exact details on the inputs and outputs of the controllers are not provided, but both policies are probabilistic in nature.

Thuruthel et al. (2019) use guided policy search (**<empty citation>**) over trajectory samples containing control actions for given regions in the state-space of a soft robotic manipulator. The resulting control policy is represented as a FNN which takes in current and previous states, as well as a desired target state, and returns the appropriate control action.

citation [17] from thuruthel2019

For the control of a soft robotic arm, in Nazeer et al. (2023), the policy is represented by a multi-layer perceptron (MLP) which takes in the arm's current state¹ and outputs a control action. The policy is trained using the Soft Actor-Critic (SAC) (**<empty citation>**) reinforcement learning algorithm.

SAC reference

In Lu et al. (2020), a more complicated technique is used. A probabilistic 'skill' distribution $p_\psi(z|s)$ and a neural network-based policy $\pi_\theta(a|s, z)$ are independently and simultaneously learned using SAC. Both models are trained on rollouts of state transitions generated by a dynamics model. However, once learned, a model-based planning algorithm called Model Path Predictive Integral (MPPI) (**<empty citation>**) is used to find the optimal skill sequence in the skill distribution p_ψ , resulting in executable actions through the policy π_θ .

citation

¹Rather, the sum of the state and a correction term, resulting from a separate network (section 4.2), but that is not relevant for the basic idea of the policy.

3.2. Learning the dynamic model

One way of having continually adapting control for soft robots, is by implementing a (traditional) control scheme that uses a dynamic model (either forward or inverse) of the robot and then to incrementally update that model. The idea of learning a dynamic model for soft robots has advantages even outside the context of lifelong adaptation: many soft robots have such a large amount of degrees of freedom and such complex and nonlinear dynamics, that analytically deriving a dynamic model is both cumbersome and often insufficiently accurate (Gillespie et al., 2018; Giorelli et al., 2015; Kim et al., 2021; Thuruthel et al., 2019).

Blackbox methods also allow for learning both forward and inverse dynamic models and can thus be used in several control schemes. When learning a forward model, the robot's actions and resulting states can be used as the inputs and outputs, respectively, this is so-called *direct modeling* (Nguyen-Tuong & Peters, 2011b). In the case of learning inverse dynamic models, such direct approaches might not always be applicable, due to the multiple solutions that exist for the inverse dynamics of robots with redundant degrees of freedom. Instead, an *indirect-modeling* approach can be taken, where the model is trained to minimize the error of the feedback controller (Nguyen-Tuong & Peters, 2011b).

The form that the dynamic model takes can vary. For example, in Gillespie et al. (2018), a neural network is trained to predict the next state of the robot, given the current state and a control action, the analytical gradients of these outputs w.r.t. to the inputs are made available during the process of backpropagation and can be used in matrices A and B to form a more interpretable state-space system. In Thuruthel et al. (2019) on the other hand, a forward model is learned for open-loop control, using the control action, the current state, and the previous state as input, and the next state as output to train a recurrent neural network. This forward model is kept in its more abstract blackbox form as a recurrent neural network (RNN), and is used in combination with trajectory optimization to generate trajectories in open-loop control.

Since the lifelong adapting nature of the techniques is embedded in the dynamic model, the choice of control scheme in which the learned dynamic model is used can vary from implementation to implementation. To illustrate, the aforementioned matrices A and B from Gillespie et al. (2018) are used in a model predictive control setting, while the generated trajectories from Thuruthel et al. (2019) are used as samples to train a blackbox control policy (the dynamic model is in essence an intermediary step in order to create the blackbox policy mentioned in section 3.1). More examples can be found in the literature.

Nazeer et al. (2023) represent a forward dynamic model of the soft robotic arm as an RNN model that uses the two preceding states and three preceding control actions to predict the next state. The model is then used with the previously discussed MLP policy.

As a result, and contrary to blackbox policy learning methods, the adaptation to new tasks is not included in the learning process of the dynamic model, but is entirely dependent on the control policy. In case control policies have been defined in advance for several tasks, these task settings might serve as a test to verify that the dynamic model generalizes well over these different settings. In situations where tasks are not known in advance, a control scheme such as an incrementally updated (inverse) dynamics model in some form of a feedback control loop might serve for simple applications (Nguyen-Tuong & Peters, 2011b). For more complex tasks, a more intricate control system might be warranted.

3.3. Incremental learning

In section 3.1 and 3.2 an overview was provided of how blackbox machine learning techniques can be used to model either control policies or dynamic models (to be used in control policies). The choice of machine learning method can be important when it comes to the desired incremental nature of the learning techniques. To illustrate, consider these two main strategies in

rethink order of paragraphs

continually adaptive control:

1. All training is performed offline (after gathering sufficient data) and the robot is now completely capable of immediately adapting to changing dynamics it will encounter during its operational lifespan.
2. Training is performed incrementally as data comes in, which can potentially be done in real-time.

In both scenarios, the control of the robot would be continually adaptive. However, only strategy 2 aligns with the incremental objectives of the adaptive control challenge. In addition, the complex dynamics and many degrees of freedom of soft robots might make it difficult to accurately represent all possible degradations in dynamics (such as damages) or state-space configurations of the robot, potentially causing an imbalance in the training dataset of strategy 1. Incrementally learning to adapt the control policy or dynamic model is thus far more desirable, since it can cope with unforeseen circumstances.

Not all the previously discussed methods are - or have the possibility to be - implemented in an incremental manner. As an example, the two policies from Zhou et al. (2022) mentioned in section 3.1 are trained sequentially: one online and the other offline. Each of these policies required different training/optimization algorithms for their specific setting, in this case Maximum a Posteriori Policy Optimization (MPO) (Abdolmaleki et al., 2018) for the online interaction setting, and Critic Regularized Regression (CRR) (Wang et al., 2020) for the offline distillation setting (see ??).

3.4. Summary and conclusions

In Table 3.1 an overview of the discussed methods can be found.

Reference	Dynamic model	Control policy	Incremental
Gillespie et al. (2018)	SS-matrices	MPC	✓
Gijssberts and Metta (2011)	Linear model in random feature space	–	✓
Thuruthel et al. (2019) ²		Open-loop + trajectory optimization	–
Lu et al. (2020)	Model ensemble	MPC	✓
Nguyen-Tuong and Peters (2011a)	Gaussian process	Feedforward	✓
Romeres et al. (2016)	Gaussian process	–	✓
Nazeer et al. (2023)	RNN	MLP	3
Thuruthel et al. (2019) ⁴	–	FNN	–
Alessi et al. (2023)	–	FNN	–
Zhou et al. (2022)	–	Probabilistic model	✓


Table 3.1: An overview blackbox methods and their ability to be implemented incrementally. For dynamic model learning methods, the control policy specifies the control scheme in which the model is used.

¹Used for generating training samples

²The incremental nature of the method is achieved by a third network, not by the dynamics model or control policy

³After training on generated samples

footnotes, checkmarks, table size, consistency: mention the form (eg. NN, gaussian process, ... or just 'blackbox'?, check two Lu2020 and Lu2019 references) or the training algorithm (eg. PPO)



Learning architectures

4.1. Dual architectures

(Lu et al., 2019, 2020; Mallya et al., 2018; Schwarz et al., 2018; Shin et al., 2017; Sprechmann et al., 2018).

4.2. Architectural expansions

(Chatzilygeroudis et al., 2016; Gillespie et al., 2018; Nazeer et al., 2023).

5

Storage and retrieval

(Aljundi et al., 2016; Mouret & Clune, 2015; Nguyen-Tuong & Peters, 2011a; Schulman et al., 2017; Shin et al., 2017; Sprechmann et al., 2018; Wang et al., 2022; Xie & Finn, 2021).

6

Discussion and conclusion

Todo list

general intro, define "blackbox"	5
citation [17] from thuruthel2019	5
SAC reference	5
citation	5
rethink order of paragraphs	6
footnotes, checkmarks, table size, consistency: mention the form (eg. NN, gaussian process, ... or just 'blackbox'?, check two Lu2020 and Lu2019 references) or the training algorithm (eg. PPO)	7

Bibliography

- Abdolmaleki, A., Springenberg, J. T., Tassa, Y., Munos, R., Heess, N., & Riedmiller, M. (2018). Maximum a posteriori policy optimisation. <http://arxiv.org/abs/1806.06920>
- Alessi, C., Hauser, H., Lucantonio, A., & Falotico, E. (2023). Learning a controller for soft robotic arms and testing its generalization to new observations, dynamics, and tasks. *2023 IEEE International Conference on Soft Robotics (RoboSoft)*, 1–7. <https://doi.org/10.1109/RoboSoft55895.2023.10121988>
- Aljundi, R., Chakravarty, P., & Tuytelaars, T. (2016). Expert gate: Lifelong learning with a network of experts. <http://arxiv.org/abs/1611.06194>
- Annaswamy, A. M., & Fradkov, A. L. (2021, January). A historical perspective of adaptive control and learning. <https://doi.org/10.1016/j.arcontrol.2021.10.014>
- Chatzilygeroudis, K., Cully, A., & Mouret, J.-B. (2016). Towards semi-episodic learning for robot damage recovery. <http://arxiv.org/abs/1610.01407>
- Gijsberts, A., & Metta, G. (2011). Incremental learning of robot dynamics using random features. *2011 IEEE International Conference on Robotics and Automation*, 951–956. <https://doi.org/10.1109/ICRA.2011.5980191>
- Gillespie, M. T., Best, C. M., Townsend, E. C., Wingate, D., & Killpack, M. D. (2018). Learning nonlinear dynamic models of soft robots for model predictive control with neural networks. *2018 IEEE International Conference on Soft Robotics (RoboSoft)*, 39–45. <https://doi.org/10.1109/ROBOSOFT.2018.8404894>
- Giorelli, M., Renda, F., Calisti, M., Arienti, A., Ferri, G., & Laschi, C. (2015). Learning the inverse kinetics of an octopus-like manipulator in three-dimensional space. *Bioinspiration & Biomimetics*, 10, 035006. <https://doi.org/10.1088/1748-3190/10/3/035006>
- Kim, D., Kim, S.-H., Kim, T., Kang, B. B., Lee, M., Park, W., Ku, S., Kim, D., Kwon, J., Lee, H., Bae, J., Park, Y.-L., Cho, K.-J., & Jo, S. (2021). Review of machine learning methods in soft robotics (G. Gu, Ed.). *PLOS ONE*, 16, e0246102. <https://doi.org/10.1371/journal.pone.0246102>
- Landau, I. D., Lozano, R., M'Saad, M., & Karimi, A. (2011). *Adaptive control*. Springer London. <https://doi.org/10.1007/978-0-85729-664-1>
- Lesort, T., Lomonaco, V., Stoian, A., Maltoni, D., Filliat, D., & Díaz-Rodríguez, N. (2020). Continual learning for robotics: Definition, framework, learning strategies, opportunities and challenges. *Information Fusion*, 58, 52–68. <https://doi.org/10.1016/j.inffus.2019.12.004>
- Lu, K., Grover, A., Abbeel, P., & Mordatch, I. (2020). Reset-free lifelong learning with skill-space planning. <http://arxiv.org/abs/2012.03548>
- Lu, K., Mordatch, I., & Abbeel, P. (2019). Adaptive online planning for continual lifelong learning. <http://arxiv.org/abs/1912.01188>
- Mallya, A., Davis, D., & Lazebnik, S. (2018). Piggyback: Adapting a single network to multiple tasks by learning to mask weights. <http://arxiv.org/abs/1801.06519>
- Mouret, J.-B., & Clune, J. (2015). Illuminating search spaces by mapping elites. <http://arxiv.org/abs/1504.04909>
- Nazeer, M. S., Bianchi, D., Campinoti, G., Laschi, C., & Falotico, E. (2023). Policy adaptation using an online regressing network in a soft robotic arm. *2023 IEEE International Con-*

- ference on Soft Robotics (RoboSoft)*, 1–7. <https://doi.org/10.1109/RoboSoft55895.2023.10121927>
- Nguyen-Tuong, D., & Peters, J. (2011a). Incremental online sparsification for model learning in real-time robot control. *Neurocomputing*, 74, 1859–1867. <https://doi.org/10.1016/j.neucom.2010.06.033>
- Nguyen-Tuong, D., & Peters, J. (2011b). Model learning for robot control: A survey. *Cognitive Processing*, 12, 319–340. <https://doi.org/10.1007/s10339-011-0404-1>
- Romeres, D., Zorzi, M., Camoriano, R., & Chiuso, A. (2016). Online semi-parametric learning for inverse dynamics modeling. <http://arxiv.org/abs/1603.05412>
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. <http://arxiv.org/abs/1707.06347>
- Schwarz, J., Luketina, J., Czarnecki, W. M., Grabska-Barwinska, A., Teh, Y. W., Pascanu, R., & Hadsell, R. (2018). Progress & compress: A scalable framework for continual learning. <http://arxiv.org/abs/1805.06370>
- Shin, H., Lee, J. K., Kim, J., & Kim, J. (2017). Continual learning with deep generative replay. <http://arxiv.org/abs/1705.08690>
- Sprechmann, P., Jayakumar, S. M., Rae, J. W., Pritzel, A., Badia, A. P., Uria, B., Vinyals, O., Hassabis, D., Pascanu, R., & Blundell, C. (2018). Memory-based parameter adaptation. <http://arxiv.org/abs/1802.10542>
- Thuruthel, T. G., Falotico, E., Renda, F., & Laschi, C. (2019). Model-based reinforcement learning for closed-loop dynamic control of soft robotic manipulators. *IEEE Transactions on Robotics*, 35, 124–134. <https://doi.org/10.1109/TRO.2018.2878318>
- Wang, Z., Chen, C., & Dong, D. (2022). Lifelong incremental reinforcement learning with online bayesian inference. *IEEE Transactions on Neural Networks and Learning Systems*, 33, 4003–4016. <https://doi.org/10.1109/TNNLS.2021.3055499>
- Wang, Z., Novikov, A., Zolna, K., Springenberg, J. T., Reed, S., Shahriari, B., Siegel, N., Merel, J., Gulcehre, C., Heess, N., & de Freitas, N. (2020). Critic regularized regression. <http://arxiv.org/abs/2006.15134>
- Xie, A., & Finn, C. (2021). Lifelong robotic reinforcement learning by retaining experiences. <http://arxiv.org/abs/2109.09180>
- Zhou, W., Bohez, S., Humplik, J., Abdolmaleki, A., Rao, D., Wulfmeier, M., Haarnoja, T., & Heess, N. (2022). Forgetting and imbalance in robot lifelong learning with off-policy data. <http://arxiv.org/abs/2204.05893>